

AI for robotics - Udacity

Localization - Monte Carlo

When you do localization using EM or DL we analyze the probability of "having x or y features in our map". We may describe it with a uniform maximum confusion function... and we update it when we explore.

Robot	□ →	7 2 3 4 5
-------	-----	-------------------

✓: clear pHit: (0-1)

✗: obstacle pMiss: (0-1)

We sense the cells and we assign prob.

to each cell. The sensing helps update the

that there is an object at any of the cells.

Motion

We may have exact or inexact motion. In exact motion we only shift the spaces or features based on the movement of the robot.

In inexact motion we assign a probability to landing in a cell that is near, but was not the target.

Having a function for sense(m_i) and one for move(m_i) entropy will decrease when we sense and increase when we move.

$$\text{Entropy} = \sum (-p \times \log(p))$$

In general, localization can be determined by the following parts:

Belief → Probability

Sense → Product, followed by normalization

Move → Convolution

Probability

$$0 \leq P(x) \leq 1, \quad \sum P(x_i) = 1$$

For measurements and where x = grid cell and

$$z = \text{measurement. } P(x_i | z) = \frac{P(z|x_i) P(x_i)}{P(z)}$$

Example Use of Bayes' Theorem (Cancer Test):

$$P(C) = 0.001$$

$$P(C|POS) = ?$$

$$P(\neg C) = 0.999$$

$$P(POS|C) = 0.8$$

$$P(POS|\neg C) = 0.1$$

Answer:

$$\bar{P}(C|POS) = 0.001 \cdot 0.8 = 0.0008$$

$$\bar{P}(\neg C|POS) = 0.999 \cdot 0.1$$

$$\alpha = 0.1007 = 0.0999$$

$$P(C|POS) = \frac{0.0008}{0.1007} = 0.0079$$

Total probability - Motion

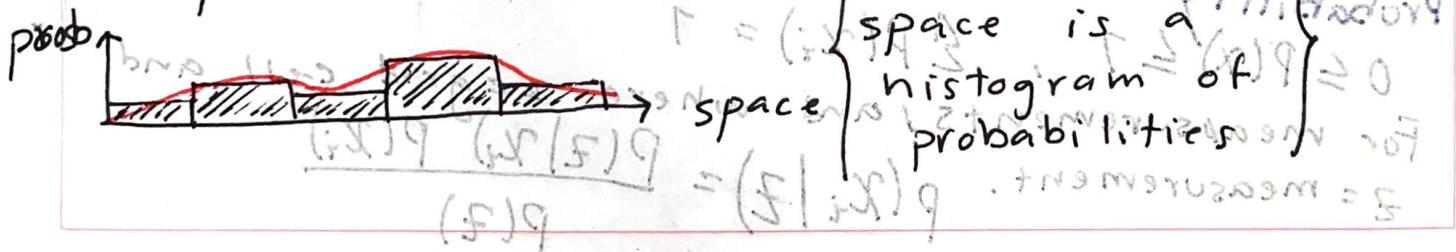
$$P(X_i^t) = \sum_j P(X_j^t) \cdot P(X_i^t | X_j) \quad \left\{ \begin{array}{l} \text{Theorem of Total Probability} \\ \text{Time} \\ \text{grid cell} \end{array} \right.$$

P_{hit} and P_{miss} are the probabilities of the sensor(s) being wrong. In the real world they are assigned by computing the noise the sensor would experience and may be set experimentally.

Tracking - Kalman Filters

They give a uni-modal, continuous approximation of the state, as opposed to the discrete multi-modal Monte Carlo localization technique.

In Monte Carlo localization the space may be as follows:

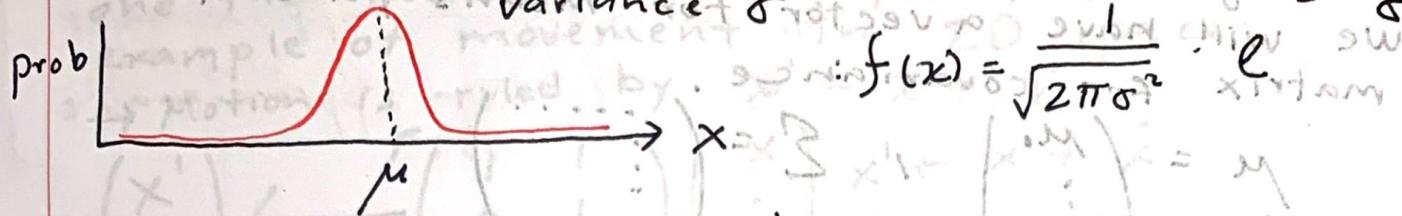


AI for robotics

D M A
written on 10/7

Scribe

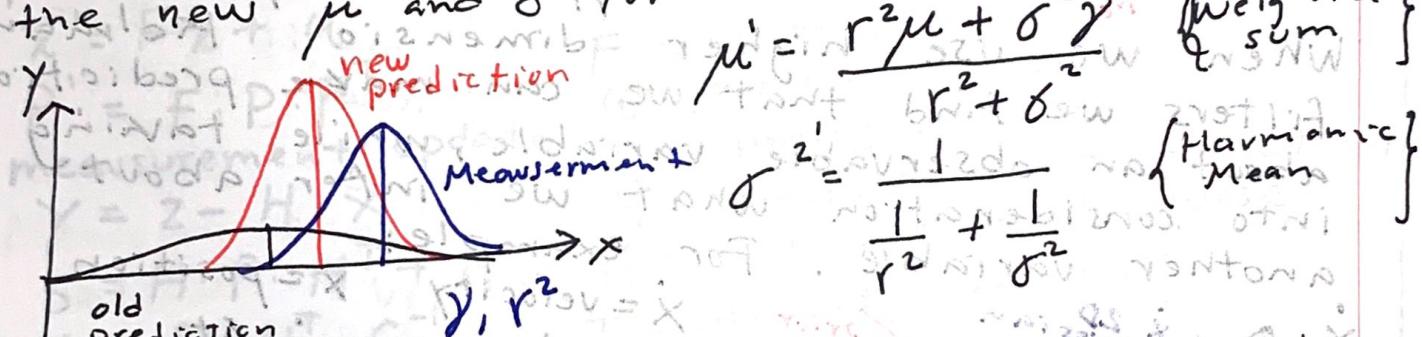
In Kalman filters we have a Gaussian distribution. Variance σ^2 is the standard deviation and prob $f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}}$



Similar to the Monte Carlo approach, in Kalman filters we have a measurement and motion update.

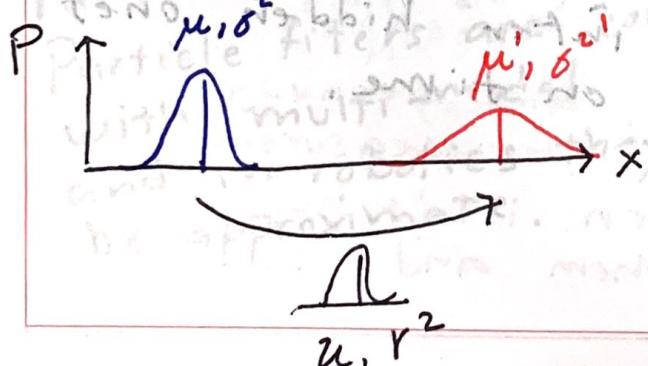
Measurement update → prediction + Total probability (motion) + convolution · product

To update our predictions we compute the new μ and σ^2 for the distribution:



The new beliefs will be more certain than either the previous belief OR the measurement.

Motion update from point A to point B
When we move from point A to point B we will have lost some certainty, due to error in movement. This is represented by:



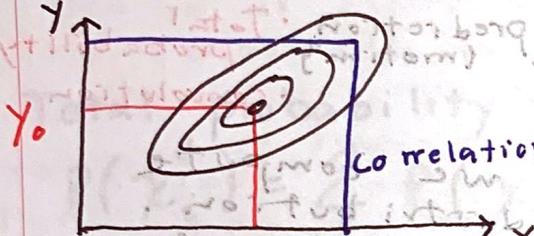
Multivariate Gaussians

In higher dimensional spaces ($S \geq \mathbb{R}^3$), we will have a vector for means, and a matrix for covariance.

$$\mu = \begin{pmatrix} \mu_1 \\ \vdots \\ \mu_n \end{pmatrix}, \Sigma = \begin{pmatrix} & & & \\ & \leftarrow & & \\ & & \downarrow & \\ & & & \rightarrow \end{pmatrix}$$

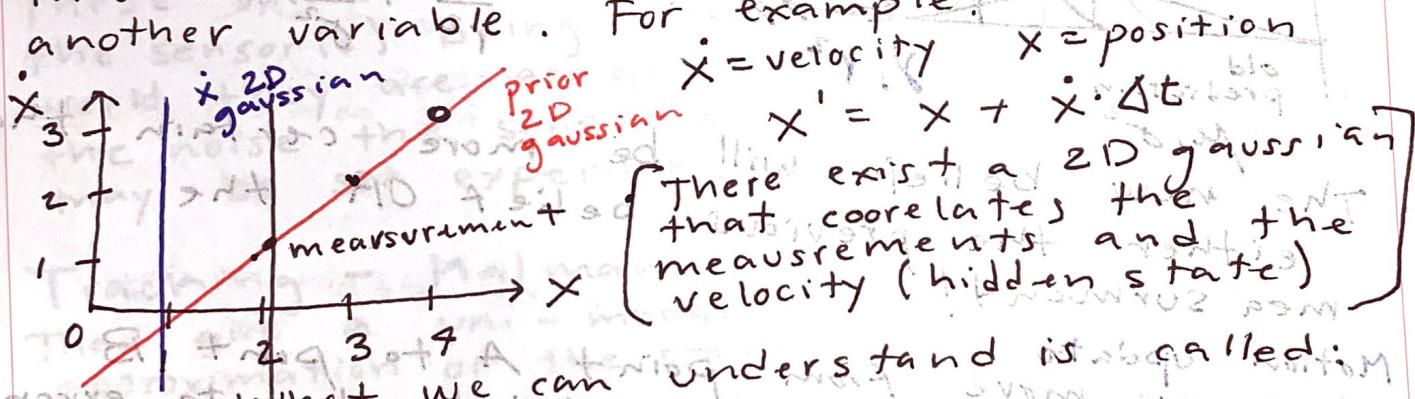
$$(S \geq \mathbb{R}^3)$$

Where the Gaussian may look like this:

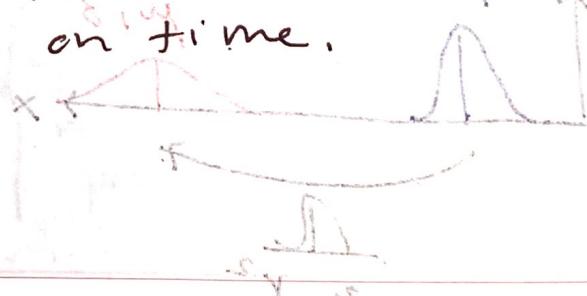


Correlation: if we move along one dimension, we would expect the other dimension to move similarly (along a curve of level of the surface or hyper-volume).

When we use higher dimensions, we find that we can make predictions about an observable variable, while taking into consideration what we infer about another variable. For example:



What we can understand is called Kalman filter states. There exist observable and hidden states. Usually, from observation of observable states, we can predict/infer hidden ones. ex. predict velocity based on time.



AI for robotics

D M A
2017-09-20 10:37

Scribe

When we design Kalman filters we need two things, one for the states and one for the measurements. For the example of movement in 2D space, motion is ruled by $\dot{x}' = x + u$

$$\begin{pmatrix} x' \\ \dot{x}' \end{pmatrix} \leftarrow \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ \dot{x} \end{pmatrix} \quad x' \leftarrow x + u$$

Where the first transformation vector H is called F and the transformation matrix P is called $F^T P F$. This is how the update is done in Kalman f.

prediction of state $x' = Fx + u$
 $F = \text{state transition matrix}$

measurement update $P' = F \cdot P \cdot F^T$
 $u = \text{motion vector}$
 $z = \text{measurement}$

$Y = z - H \cdot x$
 $H = \text{measurement function}$
 $R = \text{measurement noise}$

$S = H \cdot P \cdot H^T + R$
 $I = \text{identity matrix}$

$K = P \cdot H^T \cdot S^{-1}$

$x' = x + (K \cdot Y)$

$P' = (I - K \cdot H) \cdot P$

Notes: the "u" vector for motion is used when we may control the system, if not $u = 0$.

Particle filters are continuous state space, Particle filters are continuous belief, varying efficiency

with multimodal belief, varying efficiency and in robotics they are considered to be approximated. In robotics, particle filters are used for multi-modal belief, varying efficiency and in robotics they are considered to be approximated. In robotics, particle filters are used for multi-modal belief, varying efficiency

AI for robotics

D M A
Scriber

- Type — state space — Belief — Efficiency — In robotics
- Histogram — Discrete — Multi — exp. — approximate
- Kalman filters — continuous — Uni... — X — approximate
- Particle filters — continuous — Multi — ? belief — approximate

Particle filters

The main idea is to have particles / points placed at random in the space, and have them represent a guess of where the robot is. Particles that are more consistent with the measurements "survive" and comprise the state prediction (this surviving or not in state prediction (this comes in) is where the filter comes in).

Resampling is what is done to replace by

Resampling is what is done to replace by

$$\begin{aligned}
 & \text{We have } \sum_{i=1}^N \alpha_i = 1 \\
 & \text{and update the particles from one iteration to another. We have} \\
 & \text{weights } w_1, w_2, \dots, w_N \\
 & \text{norm. weight } \bar{w}_1 = \frac{w_1}{\sum_{i=1}^N w_i}, \bar{w}_2 = \frac{w_2}{\sum_{i=1}^N w_i}, \dots, \bar{w}_N = \frac{w_N}{\sum_{i=1}^N w_i} \\
 & \text{DRAW WITH REPLACEMENT} \Rightarrow \left[\begin{array}{c} P_1 \\ P_2 \\ P_3 \\ P_4 \\ P_5 \end{array} \right] \quad \left(\sum_{i=1}^N \bar{w}_i = 1 \right)
 \end{aligned}$$

$\alpha_1, \alpha_2, \dots, \alpha_N$

We pick points by random, and because the points that best fit the measurements have higher weights, then it's more probable that we pick them and they "survive".

continues...

AI for robotics

It is important to note that even though orientation is not checked to determine the weights of each point or particle, it does matter, since the subsequent position depend on heading.

After all, the particle filter does:

$$P(x|z) \propto P(z|x) P(x)$$

resampling importance weights \uparrow particles

Motion updates \downarrow $P(x')$

$$P(x') = \sum_{\text{samples}} P(x'|x) P(x)$$

\uparrow sample
samples

Note: A good heuristic to determine if a particle filter has enough / too many particles is to observe the weights. If they are high (non-normalized weights) you are probably tracking well and do not need many particles. The contrary is also true.

Motion planning

Search

A^* is one of the most popular and efficient search algorithms. It uses an f-value instead of the normal g-values, while still performing the same kind of search.

where $f = g + h(x, y)$ may be the function to stop, calculate the Euclidean distance to the point.